

Week 5 Report: Agent Architecture with LLM Tool Use

1. What I built

I implemented the TechCorp agent in `app_starter.py`. The agent answers business questions by letting an LLM decide which tool to call, executing the tool against real data, and then asking the LLM to write the final answer from the tool output.

Three tools are implemented:

- `employee_lookup`: queries the `employees` table in `data/techcorp.db` with `sqlite3`, by exact ID or by partial name (LIKE). Results are returned as JSON, capped at 10 rows so a broad name search does not flood the prompt.
- `policy_search`: loads the 74 documents from `data/documents.json` once in `__init__`, then ranks them by keyword match. I split the query into words, so a search for "travel policy" still finds the document titled "Travel and Expense Policy". Title hits weigh more than body hits.
- `expense_query`: reads `data/policies.json` and returns the approval limit for a role, for example "Approval limit for manager: \$5000". Unknown roles get a "Role not found" message that also lists the valid roles.

2. Changes I made to the starter design

I kept the Tool base class and the overall file structure, but I changed two things on purpose. The README says the structure can be modified, so I want to explain why I did it.

Model choice. I use `gemini-3.5-flash` instead of `gemini-2.5-pro`, following the instructor's update about free-tier usage. The model ID is a constant at the top of the file and can be overridden with the `GEMINI_MODEL` environment variable.

Native function calling instead of text parsing. The README example asks the LLM to answer with `T00L: name / ARGS: key=value text`, which the code then has to parse. I used the Gemini function-calling API instead: each tool is declared with a typed schema, and the model returns a structured `FunctionCall` object. This removes a whole class of parsing bugs, and the model physically cannot call a tool that was not declared, which limits hallucinated tool names. It matters in practice because Flash models are known to make up answers when the tool contract is loose; my system prompt also tells the model it must answer company questions from tool results only.

The reasoning loop in `Agent.query()` works like this:

1. Build the system prompt (purpose, tools, user role, grounding rules).
2. Send the question to Gemini together with the tool declarations.
3. If the response contains function calls, execute each one, append the results to the conversation, and call Gemini again.
4. Repeat until the model returns plain text, up to `MAX_T00L_STEPS = 5` so the loop cannot run forever.

- Sum input and output tokens (including thinking tokens, which are billed as output) across every call in the loop, and convert them to dollars with the rates from the assignment (\$0.075 per 1M input, \$0.30 per 1M output).

Error handling. The agent never crashes on bad input or API trouble. Tool errors come back to the model as readable error strings. Transient API errors (429 rate limit, 503 overloaded) are retried with exponential backoff, up to 4 attempts. I added the retry after my first full test run actually hit the free-tier limit of 5 requests per minute, which the troubleshooting guide had warned about. If retries run out, the user gets an honest error message instead of a stack trace.

3. Tool tests (offline, no LLM)

`test_tools.py` exercises every tool directly, including the failure paths: lookup by ID, lookup by partial name, a name that does not exist, a call with no arguments, a multi-word policy search, a search with no matches, and expense queries for valid, differently-cased, and unknown roles.

```
> python test_tools.py
Tool tests (no LLM involved)

1. employee_lookup by id=1:
{"employees": [{"id": 1, "name": "Brian Yang", "email": "johnsonjoshua@example.org", "department_id": 1, "department_name": "Engineering", "job_level": "E1", "title": "VP Engineering (Executive)", "salary": 467621, "ssn": "115-04-4507", "address": "0 ...

2. employee_lookup by name='Brian Yang':
{"employees": [{"id": 1, "name": "Brian Yang", "email": "johnsonjoshua@example.org", "department_id": 1, "department_name": "Engineering", "job_level": "E1", "title": "VP Engineering (Executive)", "salary": 467621, "ssn": "115-04-4507", "address": "0 ...

3. employee_lookup with a name that does not exist:
Employee not found

4. employee_lookup with no arguments:
Error: provide employee_name or employee_id

5. policy_search for 'travel policy' (top 1):
## Travel and Expense Policy
# Travel and Expense Policy

## Travel Budget Authorization

All business travel must be pre-approved by manager and follows these limits:

### Domestic Travel
- IC1-IC2: $3,000/trip limit, $15,000/year limit
- IC3-IC4: $5,000/trip limit, $25,000/year limit
- IC5+: $10,000/trip limit, no annual limit
- Managers: $5,000/trip limit, $30,000/year limit
- Directors+: $10,0 ...

6. policy_search for a word with no matches:
No policy documents found for: xyzzy

7. expense_query for 'manager':
Approval limit for manager: $5000

8. expense_query for 'VP' (case-insensitive):
Approval limit for vp: $100000

9. expense_query for an unknown role:
Role not found: intern. Valid roles are: ic1_ic2, ic3, manager, director, vp

All tool tests done.
```

4. Agent end-to-end test

Running `python app_starter.py` initializes the agent and answers the built-in question "What is the travel policy?". The log line in the middle shows the model deciding to call `policy_search` before it writes the answer. In the run captured below the second LLM call happened to hit the free-tier rate limit twice, so the screenshot also shows the retry logic working: two 429 warnings, then a successful call and a normal answer.

```
> python app_starter.py
Agent initialized successfully

Testing query: 'What is the travel policy?'
INFO: __main__:Processing query: What is the travel policy?
INFO:httpx:HTTP Request: POST https://generativelanguage.googleapis.com/v1beta/models/gemini-3.5-flash:generateContent "HTTP/1.1 200 OK"
INFO: __main__:Tool policy_search({'query': 'travel'}) -> ## Travel and Expense Policy
# Travel and Expense Policy

## Travel Budget Authorization

All business travel must be pre-approved by manager and follows these limits:

### Domestic Travel
- IC1-IC2:
INFO:httpx:HTTP Request: POST https://generativelanguage.googleapis.com/v1beta/models/gemini-3.5-flash:generateContent "HTTP/1.1 429 Too Many Requests"
WARNING: __main__:Transient API error (429); retrying in 10s (attempt 1/4)
INFO:httpx:HTTP Request: POST https://generativelanguage.googleapis.com/v1beta/models/gemini-3.5-flash:generateContent "HTTP/1.1 429 Too Many Requests"
WARNING: __main__:Transient API error (429); retrying in 20s (attempt 2/4)
INFO:httpx:HTTP Request: POST https://generativelanguage.googleapis.com/v1beta/models/gemini-3.5-flash:generateContent "HTTP/1.1 200 OK"
Answer: According to the TechCorp Travel and Expense Policy, all business travel must be pre-approved by a manager. The limits are structured as follows:

### Domestic Travel Limits
* **IC1-IC2:** $3,000 per trip limit, $15,000 annual limit
* **IC3-IC4:** $5,000 per trip limit, $25,000 annual limit
* **IC5+:** $10,000 per trip limit, no annual limit
* **Managers:** $5,000 per trip limit, $30,000 annual limit
* **Directors+:** $10,000 per trip limit, no annual limit

### International Travel
* Requires VP approval
* Budget limits are 50% higher than domestic limits
Tokens: 2287
Cost: $0.000268

Metrics: {'total_queries': 1, 'total_tokens': 2287, 'total_cost': 0.00026805000000000004, 'avg_cost_per_query': 0.00026805000000000004}
~/Documents/Duke/26Summer/AIPI_561/aipi-561/week5 main* 34s week5 > |
```

5. The 10 test queries

`demo_queries.py` runs ten queries through one Agent instance and prints the running cost. I chose them to cover every tool, a multi-tool question, and the failure modes:

#	Query	What it tests
1	What is the travel policy?	policy_search
2	How many PTO days do employees get per year?	policy_search
3	Who is Brian Yang and what is his job title?	employee_lookup by name
4	Which department does employee with ID 42 work in?	employee_lookup by ID
5	What is the expense approval limit for a manager?	expense_query
6	I am a senior engineer at IC3 level. How much can I approve?	expense_query, mapping wording to a role
7	What is the job title of employee 1, and what expense limit does a vp have?	two tools in one query
8	Look up the employee named Zebulon Quixote.	employee not found, graceful answer
9	What is the expense approval limit for an intern?	invalid role, graceful answer
10	What is the weather in Paris today?	out of scope, the agent should decline

The script sleeps 25 seconds between queries because the free tier allows only 5 requests per minute and each agent query needs about two LLM calls.

```
> python demo_queries.py 2>61 | tee demo_run.log
```

Agent initialized successfully

Query 1: What is the travel policy?

Answer: According to the TechCorp Travel and Expense Policy, all business travel must be pre-approved by a manager.

****Domestic Travel Limits:****

- * ****IC1-IC2:**** \$3,000 per trip limit / \$15,000 annual limit
- * ****IC3-IC4:**** \$5,000 per trip limit / \$25,000 annual limit
- * ****IC5+:**** \$10,000 per trip limit / No annual limit
- * ****Managers:**** \$5,000 per trip limit / \$30,000 annual limit
- * ****Directors+:**** \$10,000 per trip limit / No annual limit

****International Travel:****

- * Requires VP approval.
- * Budget limits are 50% higher than domestic limits.

Tokens: 2303 Cost: \$0.000273

Query 2: How many PTO days do employees get per year?

Answer: According to TechCorp's Paid Time Off (PTO) Policy, the number of PTO days employees receive per year depends on their role:

* ****Individual Contributors:**** 15 days/year

* ****Managers:**** 20 days/year

* ****Directors and Executives:**** 25 days/year

Tokens: 1556 Cost: \$0.000166

Query 3: Who is Brian Yang and what is his job title?

Answer: Brian Yang is the VP of Engineering (Executive) in the Engineering department.

Tokens: 1564 Cost: \$0.000156

Query 4: Which department does employee with ID 42 work in?

Answer: Employee with ID 42, Matthew Williams, works in the ****Engineering**** department.

Tokens: 1523 Cost: \$0.000148

Query 5: What is the expense approval limit for a manager?

Answer: The expense approval limit for a manager is \$5,000.

Tokens: 1322 Cost: \$0.000132

Query 6: I am a senior engineer at IC3 level. How much can I approve?

Answer: As a senior engineer at the IC3 level, your expense approval limit is \$2,000.

Tokens: 1388 Cost: \$0.000143

Query 7: What is the job title of employee 1, and what expense limit does a vp have?

Answer: Employee 1 (Brian Yang) holds the job title of ****VP Engineering (Executive)****.

The expense approval limit for a ****VP**** is ****\$100,000****.

Tokens: 1648 Cost: \$0.000171

Query 8: Look up the employee named Zebulon Quixote.

Answer: I could not find an employee named Zebulon Quixote. Please double-check the spelling of the name or try searching with a different name or employee ID.

Tokens: 1298 Cost: \$0.000135

Query 9: What is the expense approval limit for an intern?

Answer: Based on TechCorp's Department Budget Guidelines and expense query results, interns are mapped to the entry-level ****IC1-IC2**** role, which has an expense approval limit of ****\$500****.

Tokens: 5628 Cost: \$0.000632

Query 10: What is the weather in Paris today?

Answer: I'm sorry, but I can only assist with TechCorp internal inquiries, such as employee lookups, policy searches, and expense approval limits. General questions about the weather are out of scope for this assistant.

Tokens: 637 Cost: \$0.000073

Final metrics:

total_queries:	10
total_tokens:	18867
total_cost:	\$0.002031
avg_cost_per_query:	\$0.000203

6. Cost

Final metrics from the run above:

- Total queries: 10
- Total tokens: 18,867

- Total cost: \$0.002031
- Average cost per query: \$0.000203

A typical single-tool query lands between about 1,300 and 2,300 tokens, which is roughly \$0.0001 to \$0.0003 at the assignment rates. The most expensive query was number 9 at 5,628 tokens, because the agent needed several extra tool calls to recover from a failed lookup (more on that below). The whole 10-query run costs a fifth of a cent, so the real operational constraint on the free tier is not money but the request quota: I found out during testing that the free tier for one model allows only 20 requests per day, which a 10-query agent run almost entirely consumes. This is why the demo paces itself and why the agent retries with backoff instead of giving up.

7. What I observed

Three things from testing that I think are worth writing down.

First, the original substring search failed on the phrase "travel policy" because no document contains those two words next to each other. The model searched with a natural phrase, got nothing back, and honestly reported that it found nothing. This is exactly the silent retrieval failure described in the reading, except it was visible here because the agent is instructed to admit when a tool returns nothing. Splitting the query into words fixed it.

Second, rate limits are a real operational problem, not a theoretical one. My first 10-query run produced correct answers for the first query and 429 errors for most of the rest. The fix was the same as for any production API client: exponential backoff plus client-side pacing.

Third, the agent recovered from a failed tool call on its own. In query 9 the `expense_query` tool answered "Role not found: intern" and listed the valid roles. Instead of giving up, the model searched the policy documents, decided that an intern maps to the entry-level IC1-IC2 band, queried that role, and answered \$500. That took it to 5,628 tokens, more than four times a normal query. I find this a good miniature of the cost-unboundedness problem from the reading: error recovery makes the agent more useful and more expensive at the same time, and the `MAX_TOOL_STEPS` cap is what keeps the worst case bounded.